

Quarterly Retrieval Benchmark Report

This report summarizes the Q2 evaluation of embedding models for the internal document search system. BGE-M3 was selected as the production model based on the results below.

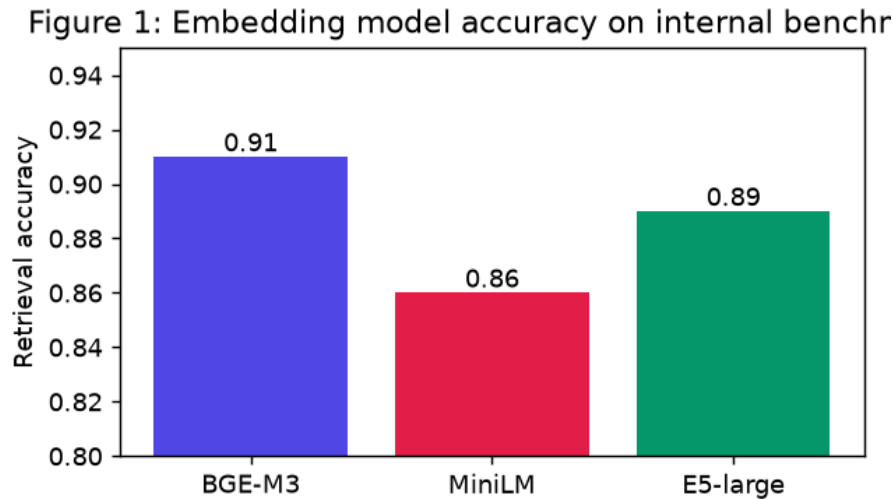


Table 1: Latency and accuracy by model

Model	Dim	Accuracy	Latency (ms)
BGE-M3	1024	0.91	42
MiniLM	384	0.86	11
E5-large	1024	0.89	58

Conclusion: BGE-M3 delivers the best accuracy (0.91) at acceptable latency and was chosen for production despite MiniLM's speed advantage.